

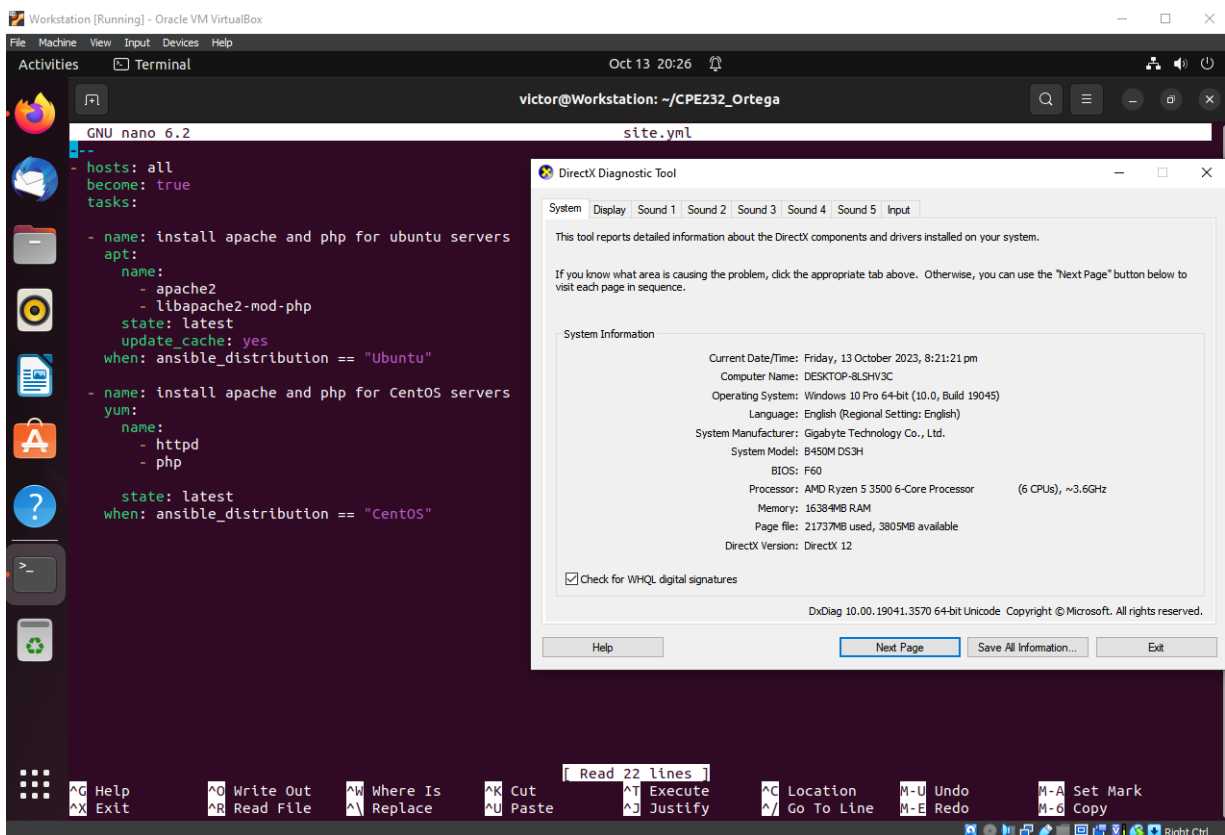
<b>Name:</b> Victor B. Ortega	<b>Date Performed:</b> 10/13/23
<b>Course/Section:</b> CPE31S5	<b>Date Submitted:</b> 10/14/23
<b>Instructor:</b> Engr. Roman Richard	<b>Semester and SY:</b> 2023-2024
<b>Activity 6: Targeting Specific Nodes and Managing Services</b>	
<b>1. Objectives:</b> 1.1 Individualize hosts 1.2 Apply tags in selecting plays to run 1.3 Managing Services from remote servers using playbooks	
<b>2. Discussion:</b>  <p>In this activity, we try to individualize hosts. For example, we don't want apache on all our servers, or maybe only one of our servers is a web server, or maybe we have different servers like database or file servers running different things on different categories of servers and that is what we are going to take a look at in this activity.</p> <p>We also try to manage services that do not automatically run using the automations in playbook. For example, when we install web servers or httpd for CentOS, we notice that the service did not start automatically.</p> <p><b>Requirement:</b>  In this activity, you will need to create another Ubuntu VM and name it Server 3. Likewise, you need to activate the second adapter to a host-only adapter after the installations. Take note of the IP address of the Server 3. Make sure to use the command <i>ssh-copy-id</i> to copy the public key to Server 3. Verify if you can successfully SSH to Server 3.</p>	
<b>Task 1: Targeting Specific Nodes</b>	

1. Create a new playbook and named it site.yml. Follow the commands as shown in the image below. Make sure to save the file and exit.

```
---
- hosts: all
  become: true
  tasks:

    - name: install apache and php for Ubuntu servers
      apt:
        name:
          - apache2
          - libapache2-mod-php
        state: latest
        update_cache: yes
      when: ansible_distribution == "Ubuntu"

    - name: install apache and php for CentOS servers
      dnf:
        name:
          - httpd
          - php
        state: latest
      when: ansible_distribution == "CentOS"
```



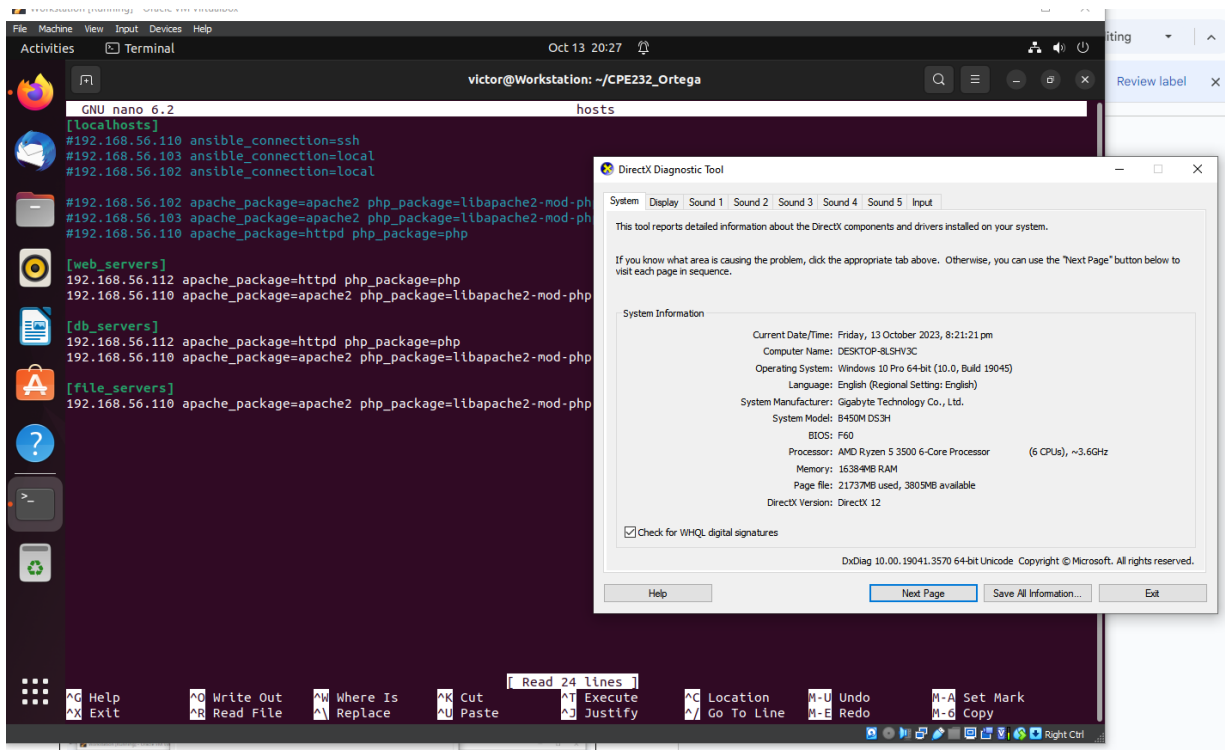
2. Edit the inventory file. Remove the variables we put in our last activity and group according to the image shown below:

```
[web_servers]
192.168.56.120
192.168.56.121

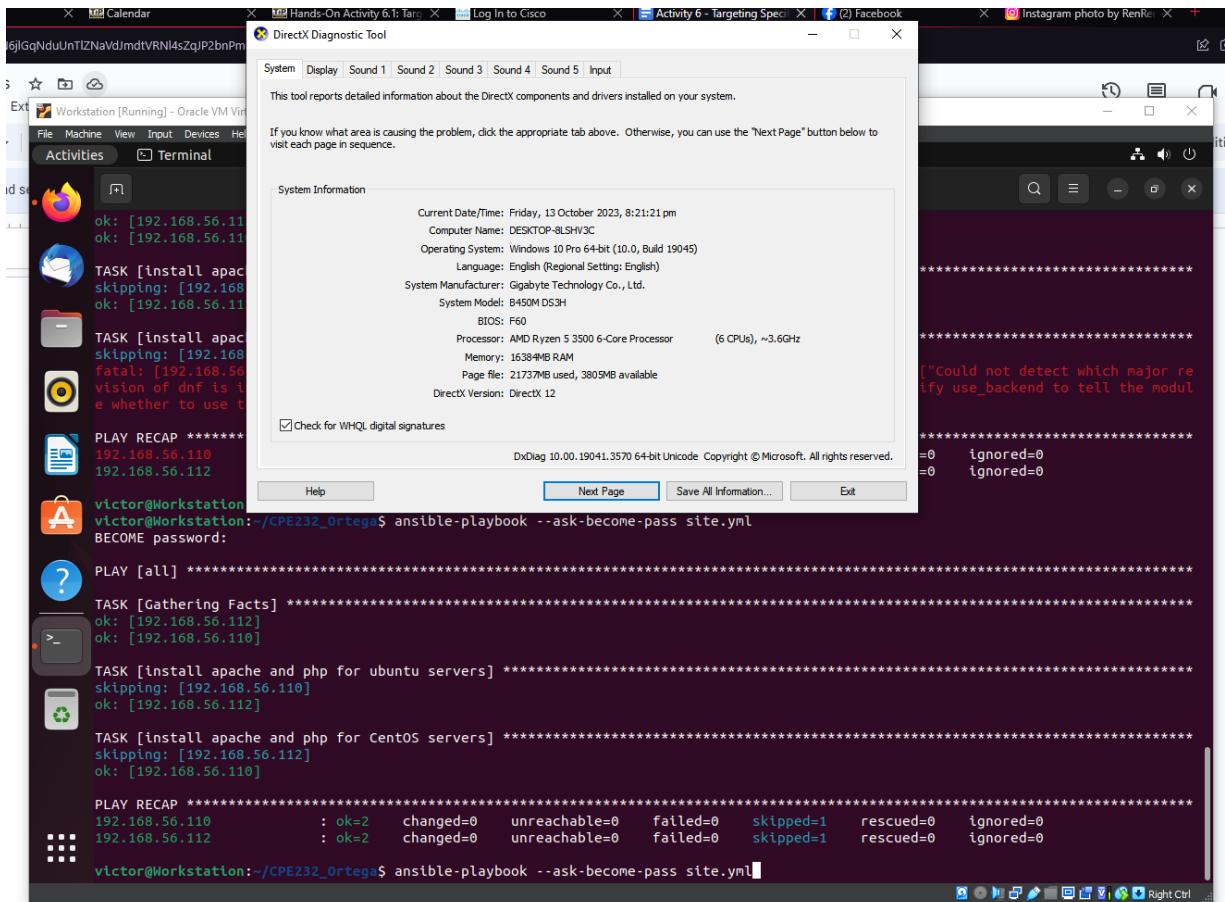
[db_servers]
192.168.56.122

[file_servers]
192.168.56.123
```

Make sure to save the file and exit.



Right now, we have created groups in our inventory file and put each server in its own group. In other cases, you can have a server be a member of multiple groups, for example you have a test server that is also a web server.



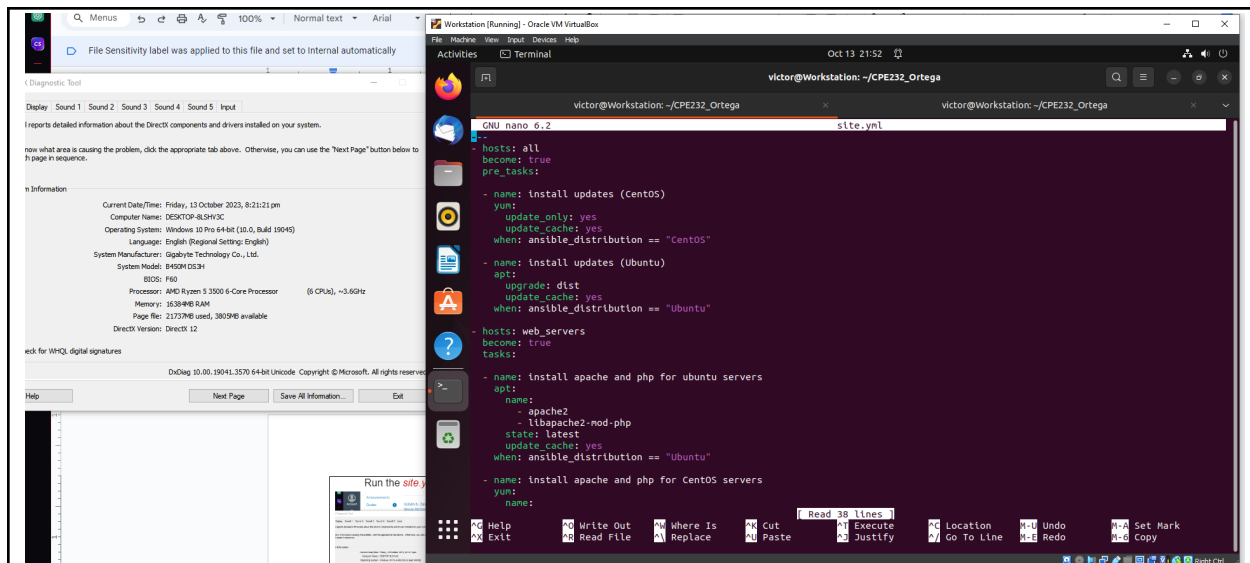
3. Edit the *site.yml* by following the image below:

```
---
- hosts: all
  become: true
  pre_tasks:
    - name: install updates (CentOS)
      dnf:
        update_only: yes
        update_cache: yes
        when: ansible_distribution == "CentOS"
    - name: install updates (Ubuntu)
      apt:
        upgrade: dist
        update_cache: yes
        when: ansible_distribution == "Ubuntu"

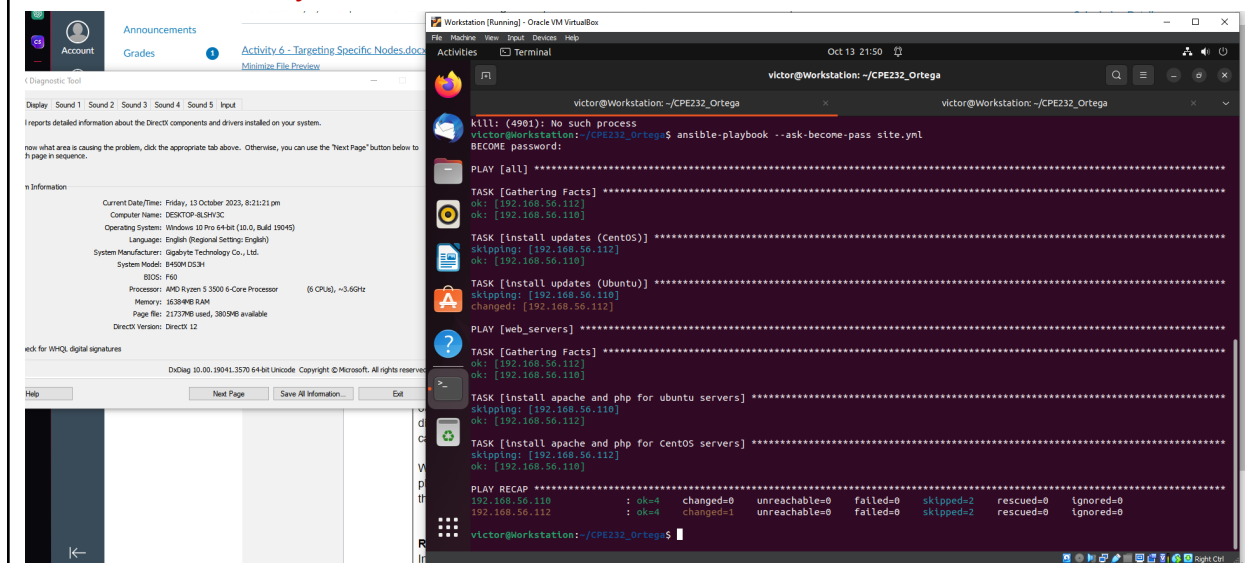
- hosts: web_servers
  become: true
  tasks:
    - name: install apache and php for Ubuntu servers
      apt:
        name:
          - apache2
          - libapache2-mod-php
        state: latest
        when: ansible_distribution == "Ubuntu"
    - name: install apache and php for CentOS servers
      dnf:
        name:
          - httpd
          - php
        state: latest
        when: ansible_distribution == "CentOS"
```

Make sure to save the file and exit.

The *pre-tasks* command tells the ansible to run it before any other thing. In the *pre-tasks*, CentOS will install updates while Ubuntu will upgrade its distribution package. This will run before running the second play, which is targeted at *web\_servers*. In the second play, apache and php will be installed on both Ubuntu servers and CentOS servers.



Run the *site.yml* file and describe the result.



4. Let's try to edit again the *site.yml* file. This time, we are going to add plays targeting the other servers. This time we target the *db\_servers* by adding it on the current *site.yml*. Below is an example: (Note add this at the end of the playbooks from task 1.3).

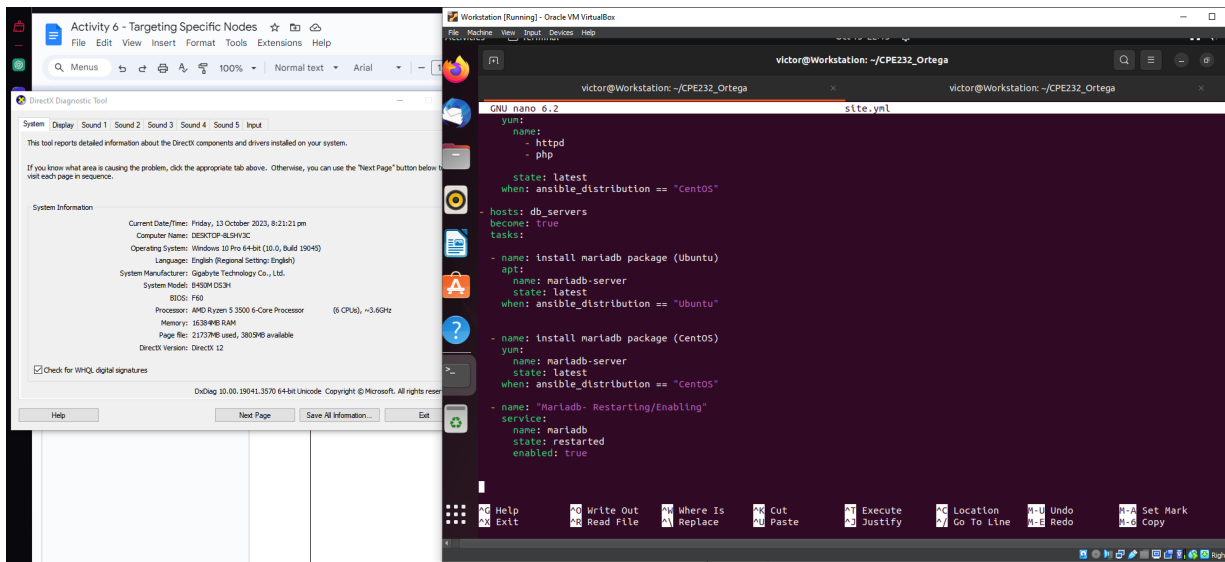
```
- hosts: db_servers
  become: true
  tasks:

- name: install mariadb package (CentOS)
  yum:
    name: mariadb-server
    state: latest
    when: ansible_distribution == "CentOS"

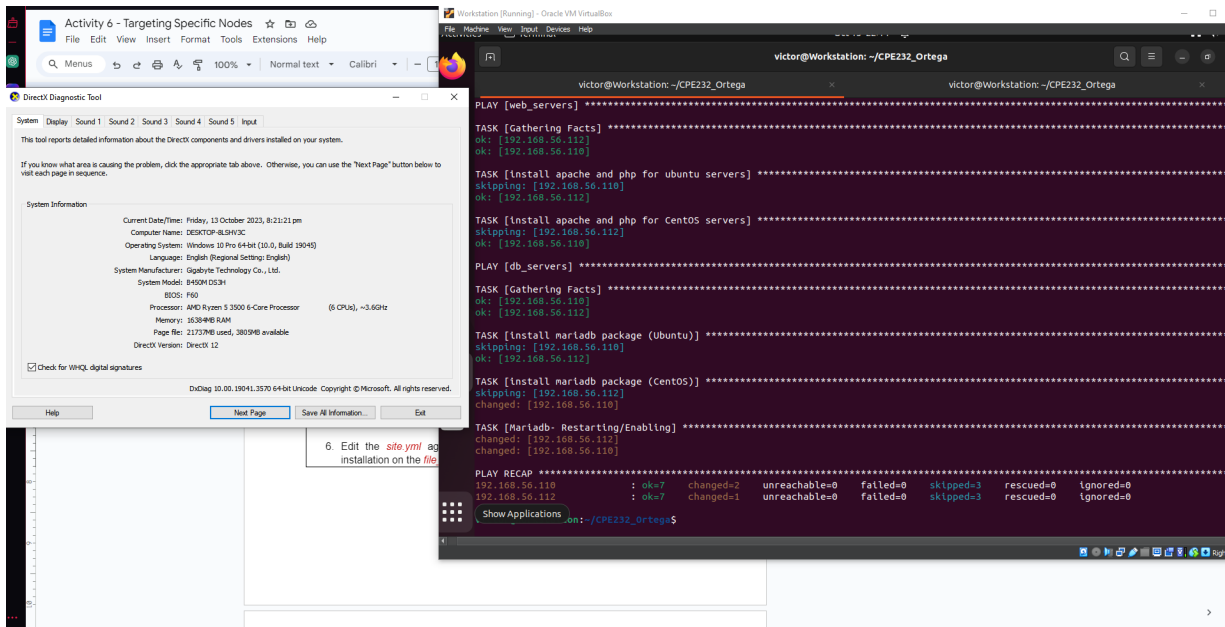
- name: "Mariadb- Restarting/Enabling"
  service:
    name: mariadb
    state: restarted
    enabled: true

- name: install mariadb package (Ubuntu)
  apt:
    name: mariadb-server
    state: latest
    when: ansible_distribution == "Ubuntu"
```

Make sure to save the file and exit.

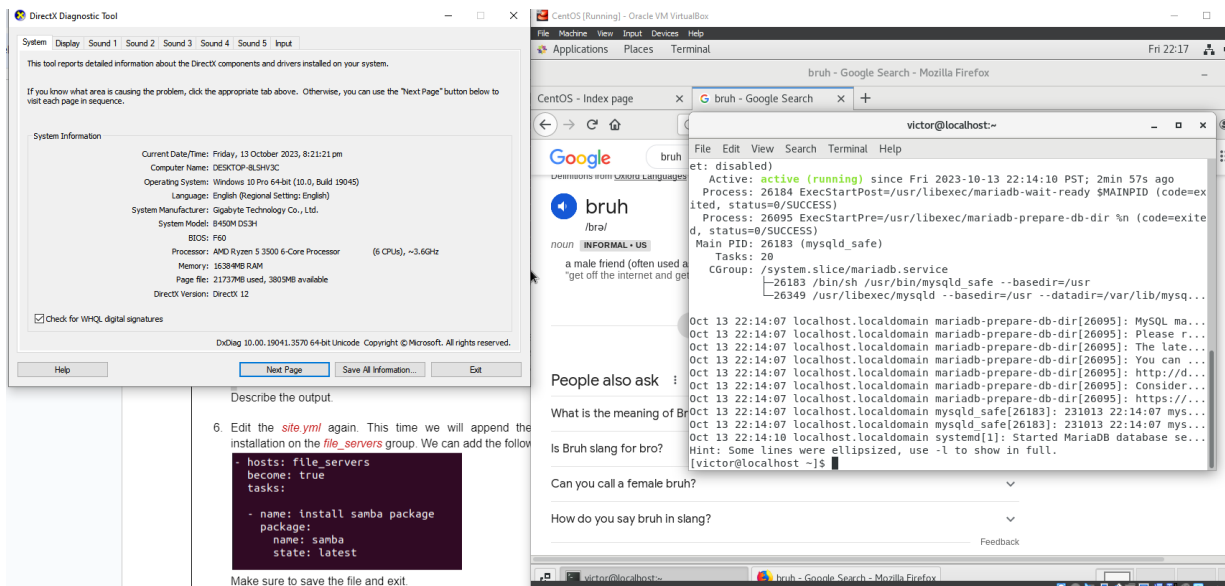


Run the *site.yml* file and describe the result.



- Go to the remote server (Ubuntu) terminal that belongs to the `db_servers` group and check the status for mariadb installation using the command: *systemctl status mariadb*. Do this on the CentOS server also.

Describe the output. The mariadb activated





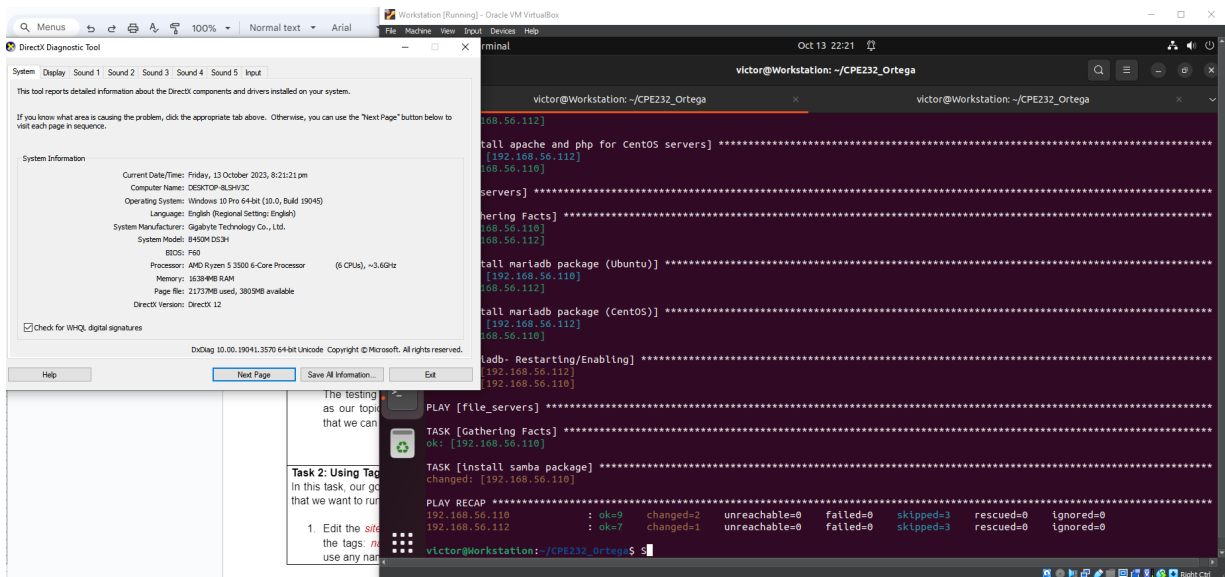
6. Edit the *site.yml* again. This time we will append the code to configure installation on the *file\_servers* group. We can add the following on our file.

```
- hosts: file_servers
  become: true
  tasks:

  - name: install samba package
    package:
      name: samba
      state: latest
```

Make sure to save the file and exit.

Run the *site.yml* file and describe the result.



The testing of the *file\_servers* is beyond the scope of this activity, and as well as our topics and objectives. However, in this activity we were able to show that we can target hosts or servers using grouping in ansible playbooks.

## Task 2: Using Tags in running playbooks

In this task, our goal is to add metadata to our plays so that we can only run the plays that we want to run, and not all the plays in our playbook.

1. Edit the *site.yml* file. Add tags to the playbook. After the name, we can place the tags: *name\_of\_tag*. This is an arbitrary command, which means you can use any name for a tag.

```

---
- hosts: all
  become: true
  pre_tasks:

    - name: install updates (CentOS)
      tags: always
      dnf:
        update_only: yes
        update_cache: yes
        when: ansible_distribution == "CentOS"

    - name: install updates (Ubuntu)
      tags: always
      apt:
        upgrade: dist
        update_cache: yes
        when: ansible_distribution == "Ubuntu"

```

```

- hosts: web_servers
  become: true
  tasks:

    - name: install apache and php for Ubuntu servers
      tags: apache,apache2,ubuntu
      apt:
        name:
          - apache2
          - libapache2-mod-php
        state: latest
        when: ansible_distribution == "Ubuntu"

    - name: install apache and php for CentOS servers
      tags: apache,centos,httpd
      dnf:
        name:
          - httpd
          - php
        state: latest
        when: ansible_distribution == "CentOS"

```

```
- hosts: db_servers
  become: true
  tasks:

    - name: install mariadb package (CentOS)
      tags: centos, db, mariadb
      dnf:
        name: mariadb-server
        state: latest
      when: ansible_distribution == "CentOS"

    - name: "Mariadb- Restarting/Enabling"
      service:
        name: mariadb
        state: restarted
        enabled: true

    - name: install mariadb package (Ubuntu)
      tags: db, mariadb, ubuntu
      apt:
        name: mariadb-server
        state: latest
      when: ansible_distribution == "Ubuntu"

- hosts: file_servers
  become: true
  tasks:

    - name: install samba package
      tags: samba
      package:
        name: samba
        state: latest
```

## Make sure to save the file and exit.

The screenshot shows two windows. On the left is the 'DirectX Diagnostic Tool' window, which displays system information. On the right is a terminal window titled 'Workstation [Running] - Oracle VM VirtualBox' showing the execution of an Ansible playbook named 'site.yml'.

**DirectX Diagnostic Tool System Information:**

- Current Date/Time: Saturday, 14 October 2023, 11:56:55 am
- Computer Name: DESKTOP-8L5H4JC
- Operating System: Windows 10 Pro 64-bit (10.0, Build 19045)
- Language: English (Regional Setting: English)
- System Manufacturer: Gigabyte Technology Co., Ltd.
- System Model: B450M DS3H
- BIOS: F60
- Processor: AMD Ryzen 5 3500 6-Core Processor (6 CPUs), ~3.6GHz
- Memory: 16384MB RAM
- Page file: 17274MB used, 7757MB available
- DirectX Version: DirectX 12

**Terminal Output (site.yml):**

```
GNU nano 6.2 site.yml
- hosts: db_servers
  become: true
  tasks:
    - name: install mariadb package (Ubuntu)
      tags: db, mariadb
      apt:
        name: mariadb-server
        state: latest
        when: ansible_distribution == "Ubuntu"
    - name: install mariadb package (centOS)
      tags: centos, db, mariadb
      yum:
        name: mariadb-server
        state: latest
        when: ansible_distribution == "CentOS"
    - name: "Mariadb- Restarting/Enabling"
      service:
        name: mariadb
        state: restarted
        enabled: true
- hosts: file_servers
  become: true
  tasks:
    - name: install samba package
      tags: samba
      package:
        name: samba
        state: latest
```

## Run the *site.yml* file and describe the result.

The screenshot shows the same two windows as before. The terminal window now displays the output of the Ansible playbook execution, showing the results of the tasks.

**Terminal Output (Results):**

```
TASK [Install apache and php for ubuntu servers] *****
skipping: [192.168.56.110]
ok: [192.168.56.112]

TASK [Install apache and php for CentOS servers] *****
skipping: [192.168.56.112]
ok: [192.168.56.110]

PLAY [db_servers] *****

TASK [Gathering Facts] *****
ok: [192.168.56.110]
ok: [192.168.56.112]

TASK [Install mariadb package (Ubuntu)] *****
skipping: [192.168.56.110]
ok: [192.168.56.112]

TASK [Install mariadb package (CentOS)] *****
skipping: [192.168.56.112]
ok: [192.168.56.110]

TASK [Mariadb- Restarting/Enabling] *****
changed: [192.168.56.112]
changed: [192.168.56.110]

PLAY [file_servers] *****

TASK [Gathering Facts] *****
ok: [192.168.56.110]

TASK [Install samba package] *****
ok: [192.168.56.110]

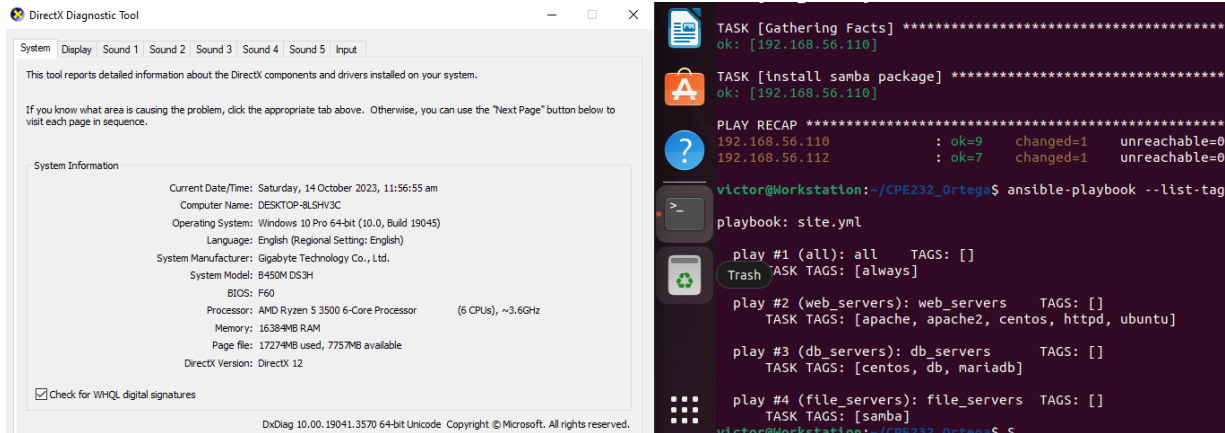
PLAY RECAP *****
192.168.56.110 : ok=9 changed=1 unreachable=0 failed=0 skipped=3 rescued=0 ignored=0
192.168.56.112 : ok=7 changed=1 unreachable=0 failed=0 skipped=3 rescued=0 ignored=0

victor@Workstation:~/CPE232_Ortega$ SSS
```

2. On the local machine, try to issue the following commands and describe each result:

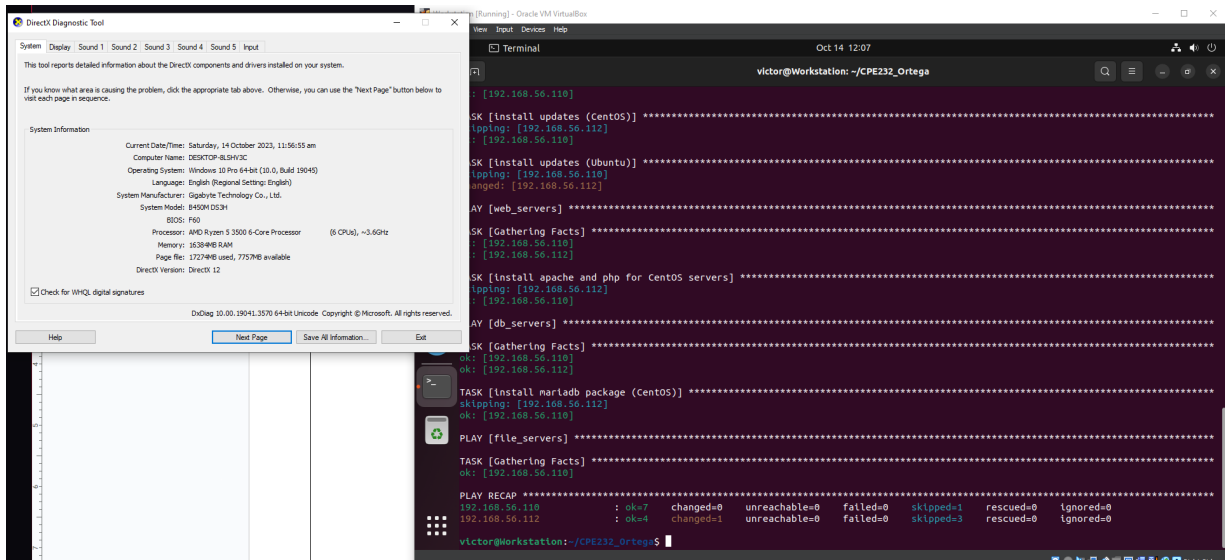
### 2.1 *ansible-playbook --list-tags site.yml*

*This command executes all commands with a tag on it.*



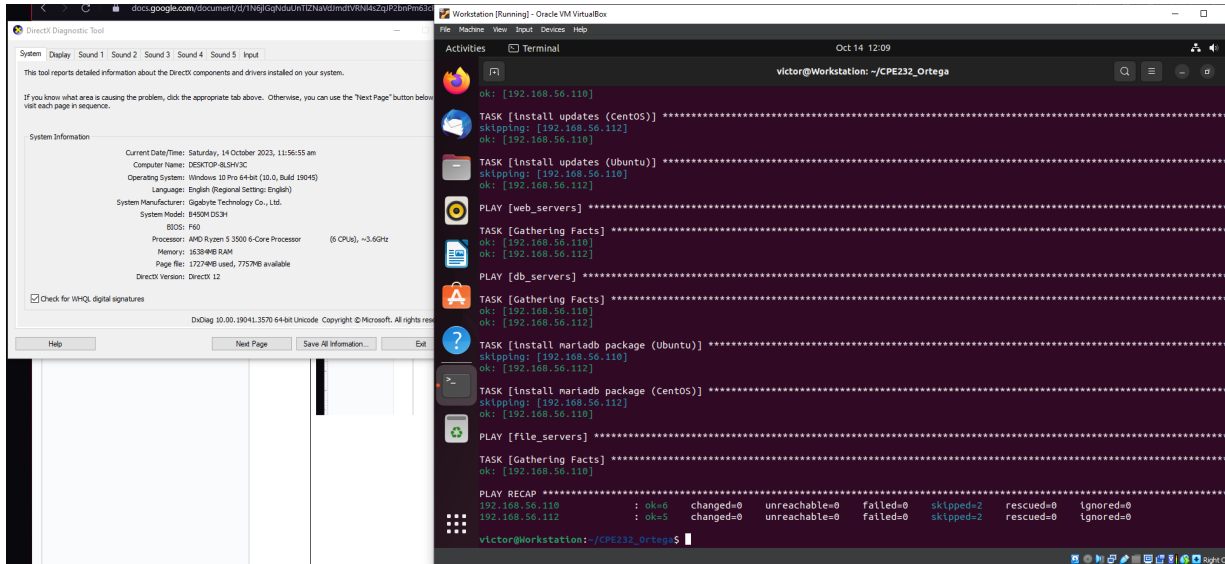
### 2.2 *ansible-playbook --tags centos --ask-become-pass site.yml*

*This command added become-pass to ask password when running the playbook*



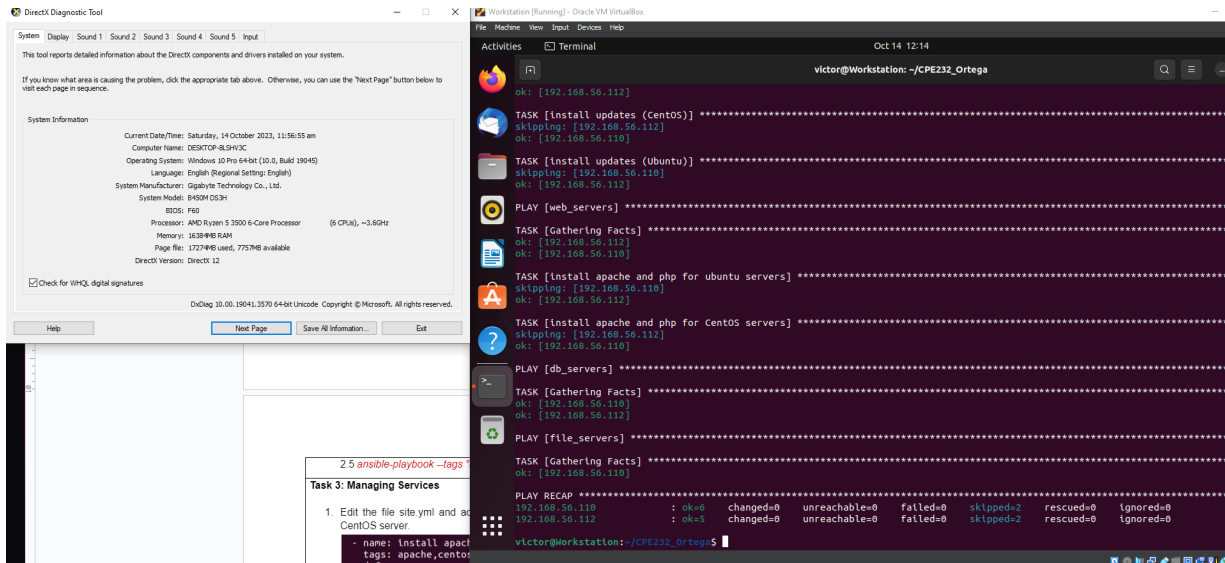
## 2.3 *ansible-playbook --tags db --ask-become-pass site.yml*

This command runs only tags with db on it.



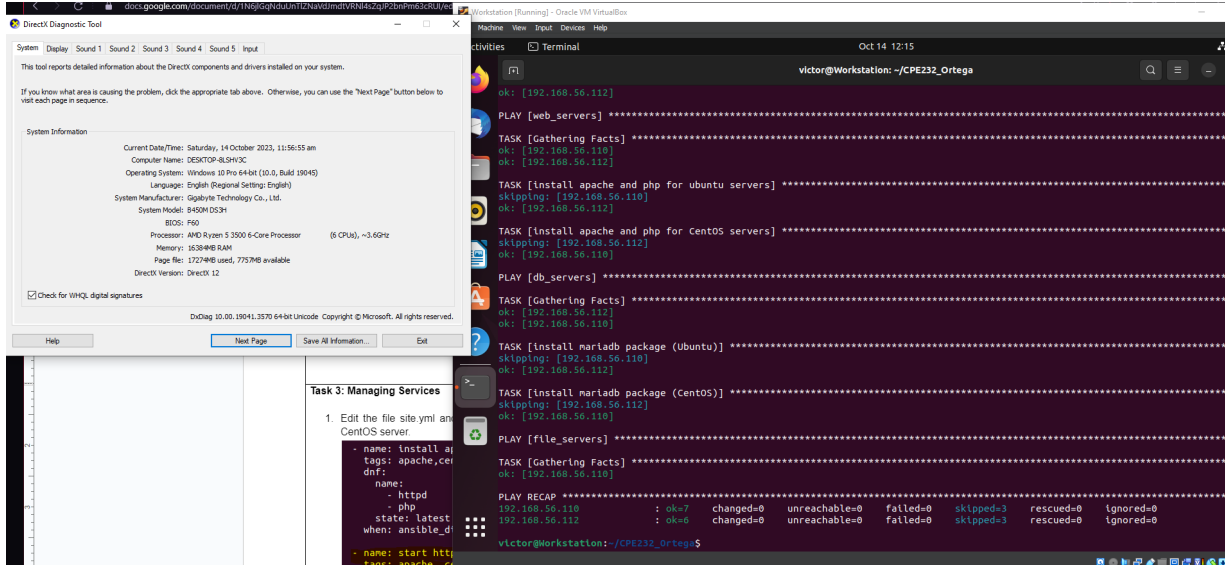
## 2.4 *ansible-playbook --tags apache --ask-become-pass site.yml*

This command runs only tags with apache on it.



## 2.5 *ansible-playbook --tags "apache,db" --ask-become-pass site.yml*

*This command runs only tags with apache and db on it.*



### Task 3: Managing Services

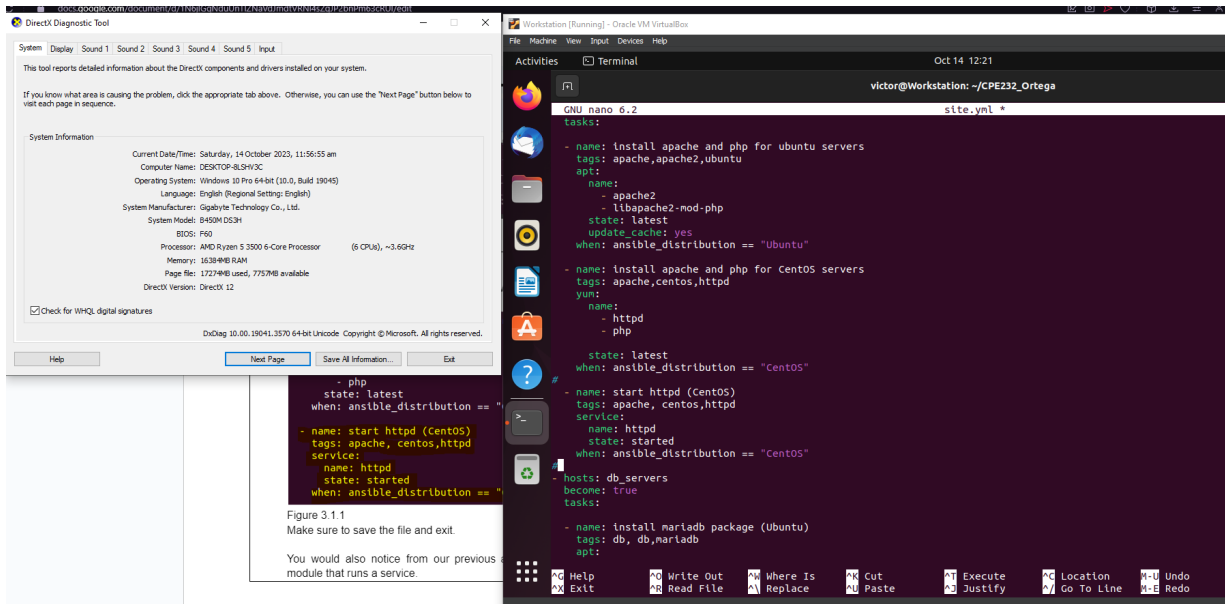
1. Edit the file site.yml and add a play that will automatically start the httpd on CentOS server.

```
- name: install apache and php for CentOS servers
  tags: apache,centos,httpd
  dnf:
    name:
      - httpd
      - php
    state: latest
    when: ansible_distribution == "CentOS"

- name: start httpd (CentOS)
  tags: apache, centos,httpd
  service:
    name: httpd
    state: started
    when: ansible_distribution == "CentOS"
```

Figure 3.1.1

Make sure to save the file and exit.



You would also notice from our previous activity that we already created a module that runs a service.

```
- hosts: db_servers
  become: true
  tasks:

    - name: install mariadb package (CentOS)
      tags: centos, db, mariadb
      dnf:
        name: mariadb-server
        state: latest
        when: ansible_distribution == "CentOS"

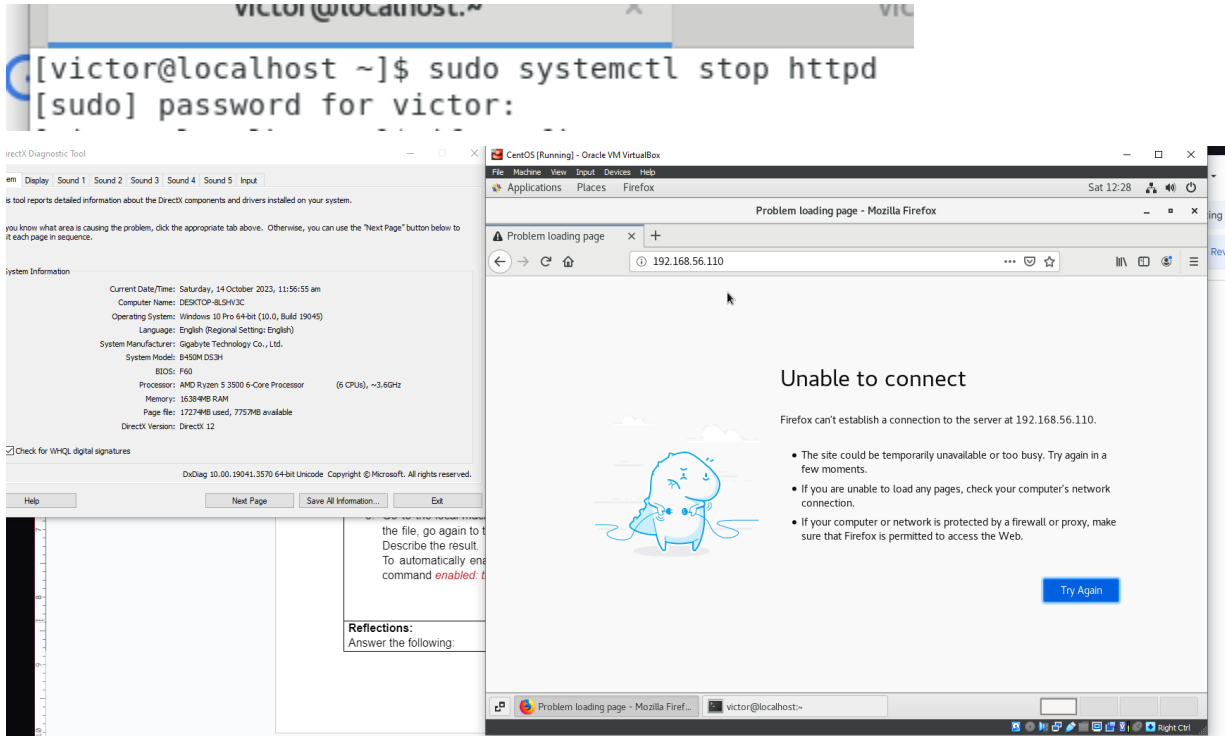
    - name: "Mariadb- Restarting/Enabling"
      service:
        name: mariadb
        state: restarted
        enabled: true
```

Figure 3.1.2

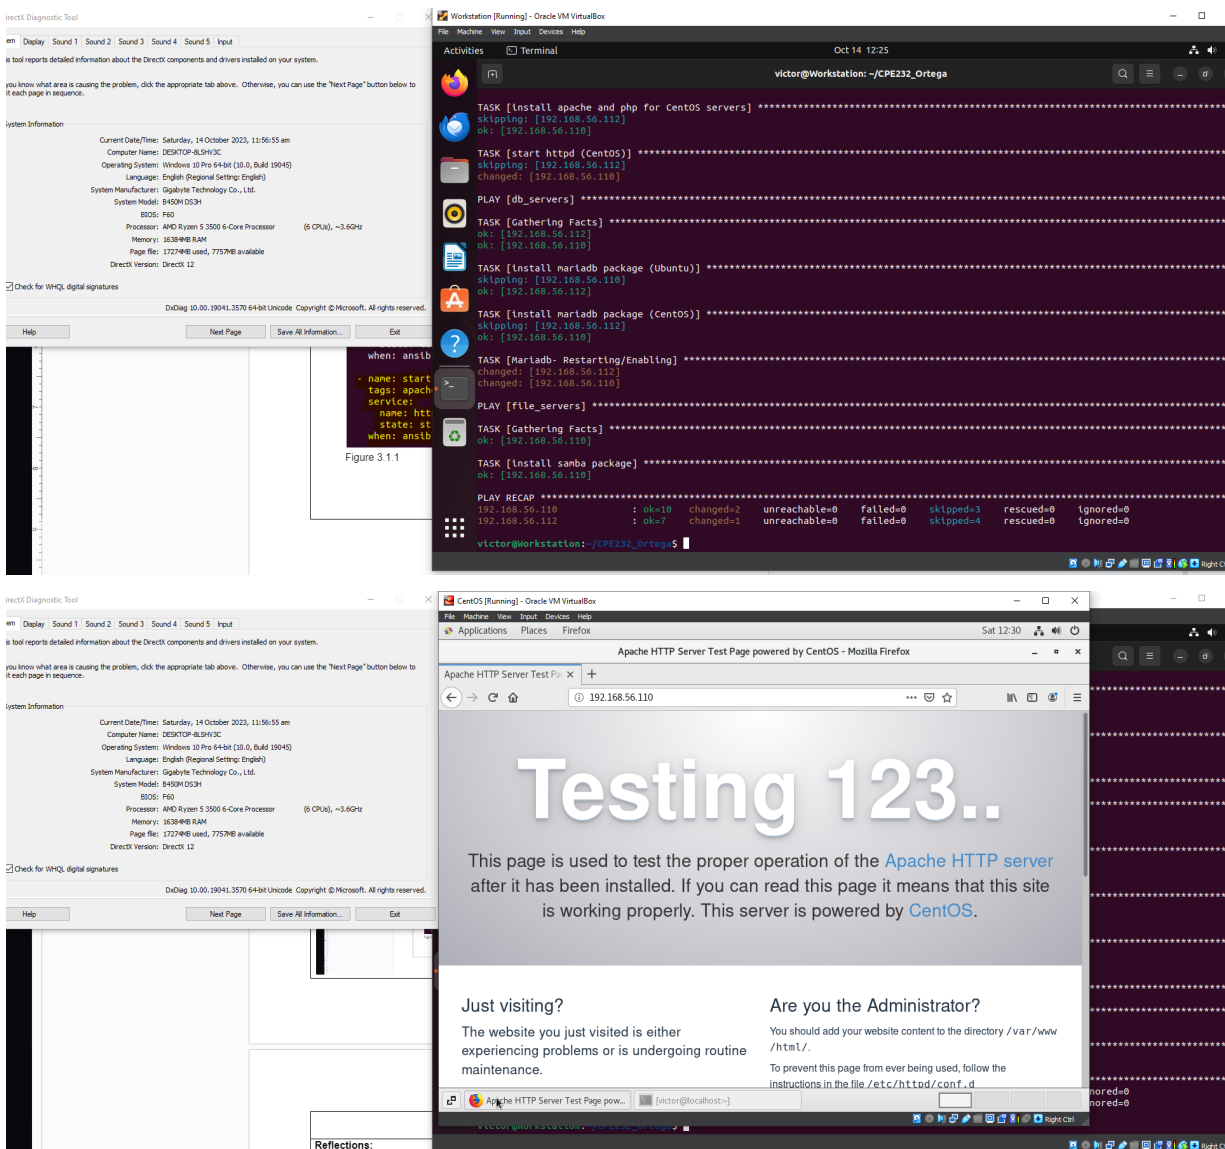
This is because in CentOS, installed packages' services are not run automatically. Thus, we need to create the module to run it automatically.



2. To test it, before you run the saved playbook, go to the CentOS server and stop the currently running httpd using the command `sudo systemctl stop httpd`. When prompted, enter the sudo password. After that, open the browser and enter the CentOS server's IP address. You should not be getting a display because we stopped the httpd service already.



3. Go to the local machine and this time, run the *site.yml* file. Then after running the file, go again to the CentOS server and enter its IP address on the browser. Describe the result.  
To automatically enable the service every time we run the playbook, use the command *enabled: true* similar to Figure 7.1.2 and save the playbook.



**Reflections:**

Answer the following:

1. What is the importance of putting our remote servers into groups?

Organizing remote servers into groups in an Ansible playbook is a powerful way to manage and execute tasks efficiently. It allows you to target specific sets of servers based on their roles, characteristics, or other criteria, and apply configurations or actions selectively. This grouping simplifies maintenance, enhances security, and facilitates scalability in large and diverse server environments.

2. What is the importance of tags in playbooks?

Tags in Ansible playbooks allow you to choose which tasks to run, which can save you time and give you more control over your configuration process. Tags make playbooks more flexible and easier to update or troubleshoot.

3. Why do think some services need to be managed automatically in playbooks?

Managing services in playbooks automates tasks, ensuring consistency, reliability, and efficiency in deploying and maintaining infrastructure. Automation reduces human error, speeds up deployment processes, and enables scalability. It also improves the ability to handle complex configurations, maintain service availability, and respond quickly to changes or issues in a dynamic environment.

**Conclusion:**

Therefore, playbooks use to automate service management makes it easier and more efficient to deploy and maintain infrastructure. It also reduces the risk of human error and improves the overall reliability and scalability of your services.

Lastly, playbooks can be used to automate a wide range of tasks, such as provisioning servers, configuring software, and deploying applications. By automating these tasks, you can free up your team to focus on more strategic initiatives.

